Outcome Ledger

Problem Validation & Feasibility Brief

Value accounting for AI-assisted engineering

Generated 2026-06-04 (UTC)

Standalone product - CPST / outcome contracts / board exports

---

Purpose of this document

Validate whether the problem is real and urgent, whether Outcome Ledger's approach (cost per accepted outcome) is feasible to build and sell, and how the current MVP proves the thesis with deterministic metrics - not adoption theater.

1. Executive summary

Enterprises are spending aggressively on AI coding tools (OpenAI, Anthropic, Cursor, Claude Code) but leadership cannot tie that spend to customer-visible outcomes. Outcome Ledger is the value accounting layer: it ingests fully loaded AI spend, defines accepted wins via versioned outcome contracts, and reports CPST (cost per stable outcome) - the unit economics finance and engineering can share.

Feasibility
graph, exe
data volum

2. Problem validation

2.1 The buyer quote (Uber COO narrative)

Andrew Macdonald (Uber COO): "That link is not there yet" - between AI adoption stats and producing more useful consumer features. Finance will compare token cost vs headcount without a line to useful features - and that trade becomes harder to justify.

2.2 What teams measure vs what leadership asks

Layer	Tracked today	Leadership asks
Inputs	Tokens, licenses, % on Copilot	-
Activity	PRs, LOC, agent runs	-
Outcomes	Fragmented	Ship velocity for users, defects, NPS
Unit economics	Rare	Cost per successful outcome by team

2.3 Failure modes (why status quo fails)

- 10x code != 10x value - review debt, wrong features, instability (DORA 2025).

- Agentic

3. Product thesis

Stop reporting AI adoption. Report what became better for the customer - and the fully loaded cost per win.

```
Workflow
  -> AI c
    -> Ou
      ->
```

4. How Outcome Ledger does costing

4.1 Total AI spend (numerator)

Sum of usage_events.cost_usd in the lookback window (default 90 days). Sources: OpenAI usage API, Anthropic admin API, Cursor/Claude Code CSV, manual CSV upload. This is fully loaded spend - not successful-call-only.

4.2 Accepted outcomes (denominator)

Counted from outcome_events matching the active outcome contract: e.g. merged PR to default branch, not reverted within stability window (OUTCOME_STABLE_DAYS). Alternative win type: default-branch commits. Only contract-defined outcome types count toward CPST.

```
Org CPST
```

4.3 Team attribution (board-ready)

```
Attributed spend % = Spend with team_id (not unassigned) / Total spend
```

Target >= 8

4.4 Outcome-linked spend (attribution graph)

For each accepted outcome, sum usage in a time window: 14 days before merge, 2 days after, preferring same GitHub repo. Produces linked %, unlinked $, avg confidence. CSV rows without repo can still link in the time window (org-wide fallback).

4.5 Weekly CPST trend

Five rolling 7-day buckets: week CPST = week spend / week stable outcomes. Green when CPST falls week-over-week (cheaper per win). Operational signal for platform leaders.

4.6 Failure cost share

Share of spend not attributable to stable outcomes or tagged as failure/retry paths - surfaces waste before finance asks.

5. Why this is valuable (not a token dashboard)

Token / usage dashboard	Outcome Ledger
We spent $X on OpenAI	We spent $X per merged PR that stuck
Adoption, seats, calls	Contract-defined accepted outcomes
Easy to game with volume	Denominator = real delivery
Finance skeptical	Versioned formula + PDF + CFO sign-off

6. Feasibility - what is built (MVP)

Capability	Status	Notes
OpenAI / Anthropic ingest	Shipped	Usage APIs
CSV usage upload	Shipped	Cursor, Claude Code
GitHub outcomes	Shipped	PR merge, revert check
Outcome contracts v1	Shipped	CFO sign-off flow
CPST + team metrics	Shipped	Deterministic
Outcome-linked graph	Shipped	Phase 1 windows
Executive report + PDF	Shipped	HITL approve
Dashboard + landing	Shipped	Railway production
Multi-tenant SaaS	Later	Design partners
Customer NPS link	Later	Phase 2

7. Feasibility scorecard

Dimension	Score (1-5)	Rationale
Pain intensity	5	Public C-suite; budget overruns
Budget availability	5	AI line item growing
Urgency (2026 proof year)	5	Investor + CFO pressure
Competition	3	Fragmented; no eng->customer winner
Willingness to pay	4	Observability + ROI SaaS comps
Ability to deliver MVP	4	Core shipped; data onboarding
Overall opportunity	Strong	Standalone category

8. Moat over time (why it compounds)

- Accepted-outcome ontology - CFO-signed definition; switching rebaselines board metrics.

- Attributio

9. Pilot checklist (prove value in 8 weeks)

- Connect OpenAI + Anthropic (or upload 90-day CSVs).

- Install G

10. Risks & honest limits

Risk	Mitigation
Outcomes company-specific	Templates + versioned contracts
Correlation != causation	Paired metrics; honest language
Sparse early data	Diagnostic ingests; CSV fallback
Competitors (Olakai, etc.)	Deeper CI/outcome chain
Privacy	No prompt storage; aggregate $ + IDs

11. Recommended next steps

- 5 CTO/CFO interviews using Macdonald quote as opener.

- Run 90-

12. References & product docs

In-repo: docs/value-one-pager.md, docs/prd.md, docs/moat.md, docs/railway-project.md. External: DORA 2025, Uber/Macdonald coverage (The Verge, Gizmodo), OptyxStack CPST methodology.