

Outcome Ledger

Problem Validation, Product Status & Pipeline June 2026

Generated 2026-06-06 07:42 UTC

What this document is for

Use this brief to validate whether the problem is real, whether Outcome Ledger solves it credibly, what has already been built, and what remains in the pipeline. Section 2 is designed for design-partner conversations and investor diligence. Section 5 lists concrete signals that prove or disprove the thesis.

Part A. Problem validation

A1. The core problem (one sentence)

Enterprises are spending millions on AI coding tools but cannot tie that spend to customer-visible outcomes. Leadership has adoption charts; finance has token bills; neither answers: what got better for users, and what did each real win cost?

A2. Anchor evidence: Uber COO (May 2026)

Andrew Macdonald (Uber President & COO) on the Rapid Response podcast and in press (The Verge, Quartz, TechSpot, May 2026):

"That link
useful con

A3. What organizations track vs what leadership asks

Layer	Tracked today	Leadership asks
Inputs	Tokens, licenses, % on Copilot	(not asked directly)
Activity	PRs, LOC, agent runs, commits	(proxy metrics)
Outcomes	Fragmented across tools	Ship velocity for users, defects, NPS
Unit economics	Rare	Cost per successful outcome by team/workflow

A4. Why the status quo fails (four failure modes)

* 10x code is not 10x value: review debt, wrong features, instability (DORA 2025).

* Agenti

A5. Market demand signals (2025-2026)

Signal	Source / implication
88% plan to increase GenAI spend	Wharton: budget is committed
61% more pressure to prove AI ROI	Kyndryl 2025: urgency rising
Only 14% CFOs see substantial ROI today	RGP 2026: expectations gap
CPST emerging as industry metric	OptyxStack, Fireworks, FinOps blogs
2026 = show me the money year	Investor + board pressure on AI line items

A6. Competitive white space

Category	Examples	Gap
LLM observability	Langfuse, Helicone	Cost per trace, not per accepted win
AI FinOps	StackSpend, Finout	Budgets/chargeback, not customer outcomes
ROI SaaS	Olakai, Pay-i, Roiva	Often surveys/hours saved, thin CI link
Spend governance	Agent wallets	Authorization, not value proof
Outcome Ledger	(this product)	Spend + GitHub win + signed definition

A7. Problem validation scorecard

Dimension	Score 1-5	Evidence
Pain intensity	5	Public C-suite quotes; budget overruns
Budget availability	5	AI line item growing double digits
Urgency	5	2026 proof year; investor pressure
Problem is widespread	4	Uber is loudest; pattern fits large tech
Willingness to pay	4	Observability + ROI SaaS comps exist
Category clarity	3	Buyers do not search CPST yet; education cost
Overall	Strong	Problem validated; product category forming

A8. Validation interview script (5 conversations)

Goal: confirm the problem exists, learn what counts as a win, and test willingness to pay for a value layer. Do not pitch features first. Lead with the Uber quote.

"Uber's COO said they cannot link AI coding spend to useful consumer features. Is that showing up for you too?"

* What do you report to the board on AI engineering spend today?

* They already built a spreadsheet or hired a consultant for AI ROI.

* They only care about inference cost per request (observability buyer).

A9. 8-week design partner proof plan

* Week 1: Connect OpenAI + Anthropic + GitHub; run first Sync.

Opening

Discover

* What v

* CFO o

* No Gith

* Week 2

Part B. Product thesis and metric

B1. What Outcome Ledger is

Outcome Ledger is the value accounting layer between AI spend and real wins. It is standalone (not Authon, not Agent Money). It does not block spend. It produces a defensible cost per win number leadership can share.

Cost per
Numerator
Denominator

B2. Product stack (six layers)

Layer	What it does
1. Vendor ingest	OpenAI, Anthropic, CSV, Langfuse traces
2. Accepted outcomes	GitHub PRs/commits, revert scan, stability gate
3. Attribution graph	Links spend to outcomes with confidence scores
4. CPST engine	Org, team, workflow CPST + monthly snapshots
5. Outcome contracts	Versioned win definition + CFO sign-off
6. Executive reports	PDF/CSV board pack; LLM narrative from metrics onl

B3. Production URLs (June 2026)

* API: outcome-ledger-production.up.railway.app

* Dashbo

Part C. What is built (shipped in detail)

C1. Core platform (MVP)

Area	Capability	Status
Ingest	OpenAI usage API	Shipped
Ingest	Anthropic admin cost API	Shipped
Ingest	CSV upload (Cursor, Claude Code)	Shipped
Ingest	GitHub merged PRs + revert check	Shipped
Ingest	GitHub default-branch commits (alt win)	Shipped
Ingest	Langfuse public API traces	Shipped (Phase 2)
Auth	Clerk multi-tenant + API keys	Shipped
Auth	MCP private agent ingest	Shipped
Metrics	Overview, teams, CPST history	Shipped
Metrics	Wins panel per PR	Shipped
Contracts	Outcome contracts draft/publish/CFO sign	Shipped
Reports	CSV export, PDF board pack	Shipped
Reports	Executive report + human approve	Shipped
Ops	Sync pipeline + audit log + daily cron	Shipped
Ops	Railway monorepo deploy (api/dashboard/landing)	Shipped

C2. Intelligence Phase 1 (v0.3)

* attribution_links table: persisted usage <-> outcome graph

* Proportion

C3. Intelligence Phase 2 (v0.4)

* Learned linker: logistic regression on overrides + high-confidence links

* Usage

C4. Dashboard and landing (latest)

* Overview: CPST chart, teams table, benchmark intelligence, outcome graph

* Integration

C5. Key API endpoints

Endpoint	Purpose
----------	---------

POST /v1/sync	Full ingest + graph rebuild + snapshots
GET /v1/metrics/overview	Org CPST, spend trend, teams
GET /v1/metrics/benchmarks	Deltas, anomalies, network percentiles
GET /v1/metrics/attribution	Linked spend %, sample links
POST /v1/attribution/rebuild	Rebuild attribution graph
GET /v1/attribution/candidates	Low-confidence links for review
GET /v1/reports/export.pdf	Board-ready PDF
POST /v1/reports/executive	Generate exec narrative

C6. What is NOT yet proven (honest gaps)

* No paying design partner logos with published CPST outcomes yet.

* Network

Part D. Pipeline (what comes next)

D1. Intelligence Phase 3

- * Linear/Jira linker: ticket <-> PR <-> trace graph

* XGBoost

D2. Enterprise roadmap (from docs)

Phase B - Enterprise connect (next)

- * Platform-hosted GitHub OAuth for all customers

* GitHub

- * Multi-tenant org provisioning at scale

* GitHub

D3. PRD Phase 2-3 (product roadmap)

- * Phase 2: Langfuse OTEL (partially shipped), workflow LLM classifier, Linear/Jira

* Phase

D4. Go-to-market pipeline

- * Phase 1 GTM: 3-5 design partners (Uber-shaped: large eng, coding agents, ROI pressure)

* Offer: f

D5. Recommended validation next steps (priority order)

- * 1. Run 5 CTO/CFO/COO interviews using script in Section A8.

* 2. Lan

Part E. How to use this for problem validation

E1. Hypothesis to test

Hypothesis: Engineering leaders with >\$500k/year AI coding spend cannot defend that spend to finance without a cost-per-accepted-outcome metric tied to GitHub wins.

E2. Minimum evidence to call problem validated

- * 3+ interviews where buyer describes same gap unprompted (after Uber quote).
- * 1 design

E3. Minimum evidence to call solution validated

- * Attributed spend >=80% for a real customer org.
- * Outcome

E4. Risks to disclose in validation conversations

Risk	Honest answer
Correlation not causation	We report paired metrics; we do not claim AI cause
Win definition varies	Versioned contracts; customer defines acceptance
Early data sparse	Structure works; magnitude needs real vendor inges
Competition	Observability adjacent; we own outcome-linked unit

Appendix. References

- * In-repo: docs/value-one-pager.md, docs/prd.md, docs/moat.md
- * In-repo